

David R. Treadwell IV

Yunxuan Rao

Daniel Bi

Bruce A. Maxwell

Northeastern University

## Abstract

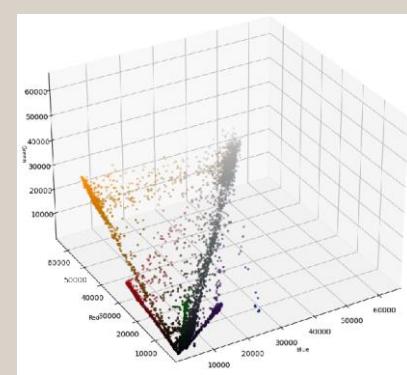
We present a lightweight fully convolutional network for color constancy (LHCC). The network uses multiple 2-D projections of the 3-D log RGB histogram of an image in order to predict the color correction coefficients.

In developing the network, we explored whether to use linear RGB or log RGB data, the network structure (width and depth), how to handle dark pixels, how to generate the 2-D mappings of the 3-D histogram, and how to normalize or transform the bin counts in order to preserve the fine histogram structure. Our results show that attention to each of these details makes a difference in overall performance.

## 3-D Histograms Represent the Colors in an Image



Original Image



Linear RGB histogram

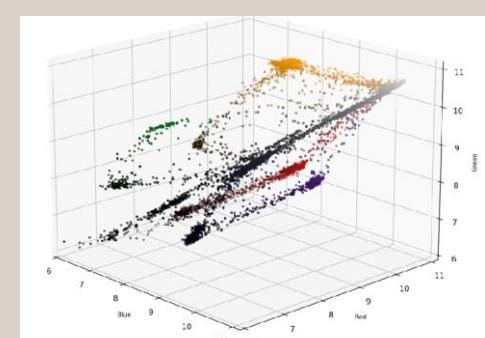
**Gamut-based methods (including DNNs) use the color distribution to estimate the most likely illumination color**

**Image Model:**  
 $I = \text{Image}$   
 $R = \text{Body Reflectance}$   
 $L = \text{Direct Illuminant}$   
 $A = \text{Ambient Illuminant}$   
 $\mathbf{l} = \text{Strength of } L_2 [0, 1]$

$$I = AR + LR = R(A + L)$$

## Log RGB separates reflectance and illumination

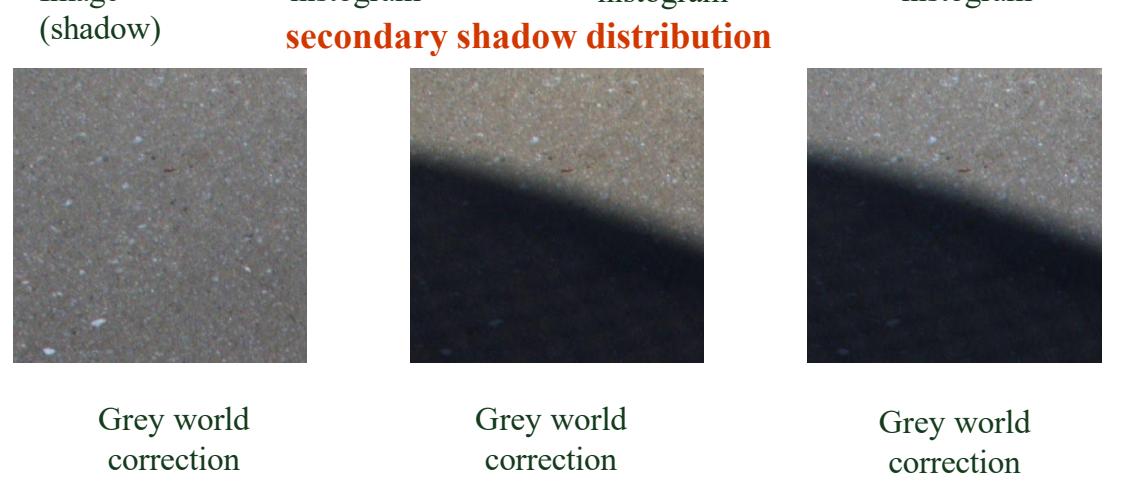
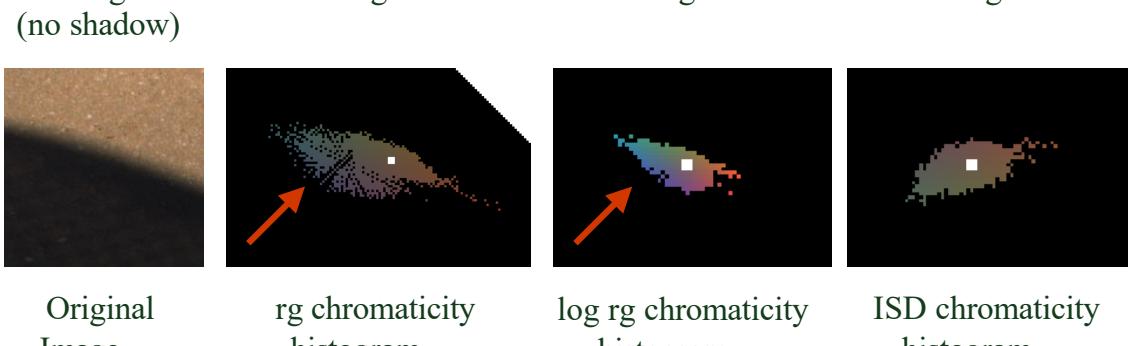
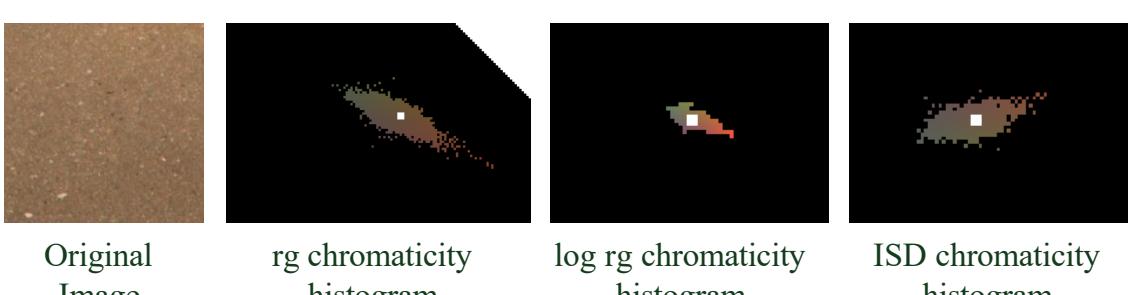
$$\log I = \log [R(A + , L)] = \log R + \log(A + , L)$$



Log RGB histogram

Note consistent lines with clear separation for each material

**Differing ambient and direct illuminants create separate, overlapping distributions of colors**



## Efficiently Analyzing a 3-D Distribution

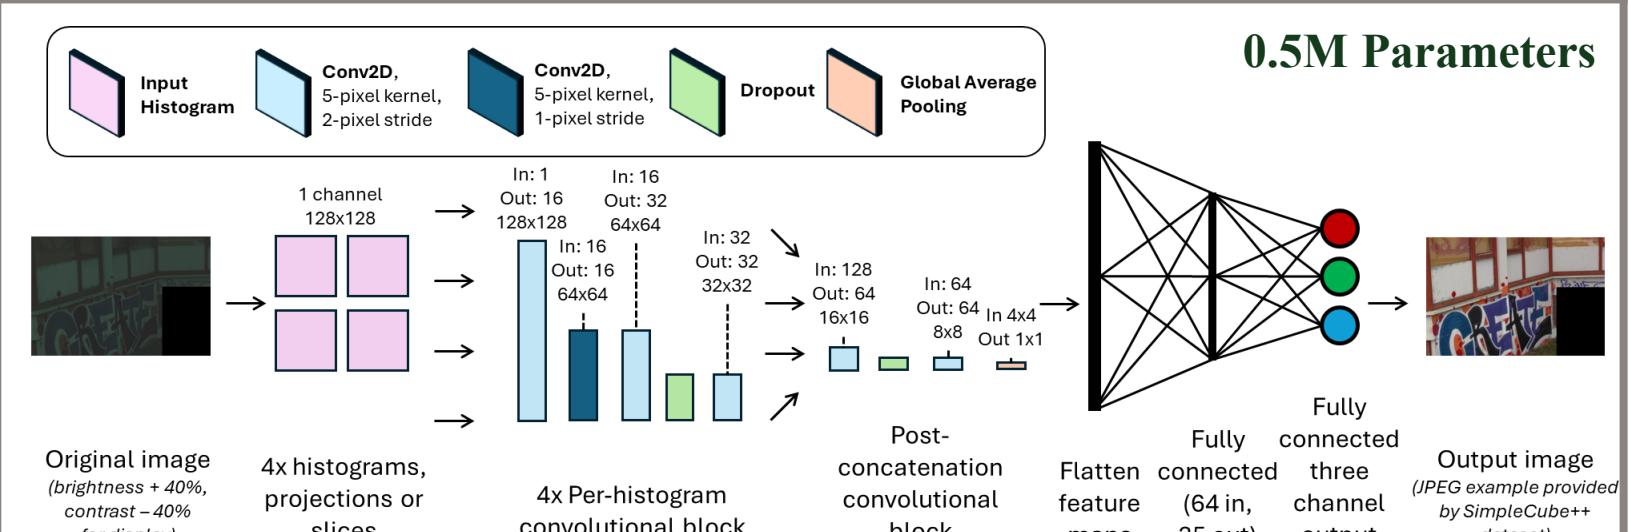
- 3-D convolutions are too computationally complex for color constancy
- Project the 3-D data onto multiple planes
- Independently extract features and then combine

## Implementation Details

- What color space to use: RGB or log RGB
- Pre-processing log space: clip to [1, 255] or [257, 65535] before taking the log
- What planes to use: {RG, RB, GB, (1, 1, 1)} or 4 slices along the 2nd eigenvector
- How to present the histogram values: raw, thresholding, hyperbolic tangent, or log
- How to design the network to integrate the separate plane information

## Exploration Results

- Color space: log RGB outperforms linear RGB for this task and network
- Planes: planar projections slightly outperform eigenvector slices
- Histogram counts: All saturation methods outperformed raw, log was the best
- Network design: slow fusion with 4 convolutions before and 2 after concatenation



Method	Mean	Med	Tri	25% Best	25% Worst	Parameters
Gray World [31]	3.18	2.00	2.37	-	-	-
Grey World 2OE [31]	3.06	1.75	2.08	-	-	-
FFCC (train/test)	2.64	1.75	1.85	0.52	6.35	12,288
FFCC (3-fold)	1.26	0.59	0.69	<b>0.18</b>	3.55	12,288
FC4 (log) [14] [24]	2.83	1.17	1.51	0.35	8.31	1,705,284
FC4 (default) [14] [24]	1.65	1.03	1.15	0.32	4.02	1,705,284
Afifi & Brown (SIIIE, Cube) [1]	1.98	1.36	-	0.40	4.64	1,008,044
Afifi & Brown (SIIIE, Cube+) [1]	2.14	1.44	-	0.44	5.06	1,008,044
CCMNet (Cube+) [16]	1.68	1.16	1.26	0.38	3.89	-
CauNet (Cube+) [18]	1.61	0.97	1.08	-	4.08	-
GoogLeNet (linear, ours)	1.53	0.95	1.05	0.25	3.90	6,797,700
GoogLeNet (log, ours)	1.35	0.75	0.88	0.22	3.52	6,797,700
LHCC (log, slices, ours)	1.12	0.63	0.69	0.19	2.93	490,895
LHCC (linear, proj, ours)	1.11	0.60	0.66	<b>0.18</b>	2.99	490,895
LHCC (log, proj, ours)	<b>1.03</b>	<b>0.57</b>	<b>0.64</b>	<b>0.18</b>	<b>2.72</b>	490,895

Method	Mean	Med	Tri	25% Best	25% Worst	Parameters
Grey World [2]	4.59	3.46	3.81	1.16	9.85	-
Grey World 2OE [2]	3.36	2.70	2.80	0.89	7.14	-
CCMNet [16]	2.32	1.71	1.83	0.53	5.18	-
FC4 (SqueezeNet-FC) [14]	2.23	1.57	1.72	0.47	5.15	1,705,284
Afifi & Brown (SqNet) [1]	2.05	1.50	-	0.52	4.48	1,008,044
FFCC [2]	1.99	1.31	1.43	<b>0.35</b>	4.75	61,480
C4 (SqNet) [32]	1.96	1.42	1.53	0.48	4.40	1.2M+
TLCC [28]	<b>1.61</b>	<b>1.27</b>	<b>1.33</b>	0.44	<b>3.35</b>	6,406,549
LHCC (pretrained, linear, ours)	2.09	1.51	1.63	0.49	4.68	490,895
LHCC (pretrained, log, ours)	2.05	1.46	1.58	0.46	4.64	490,895

## Conclusions

- Training on log RGB data consistently improves performance on all data sets
- Normalizing histogram bin counts using a log transform improves performance
- LHCC is a light-weight 0.5M parameter network robust to overtraining
- LHCC gives a new state-of-the-art on SimpleCube++ with a 2.72° error on the worst 25%